

Goal:

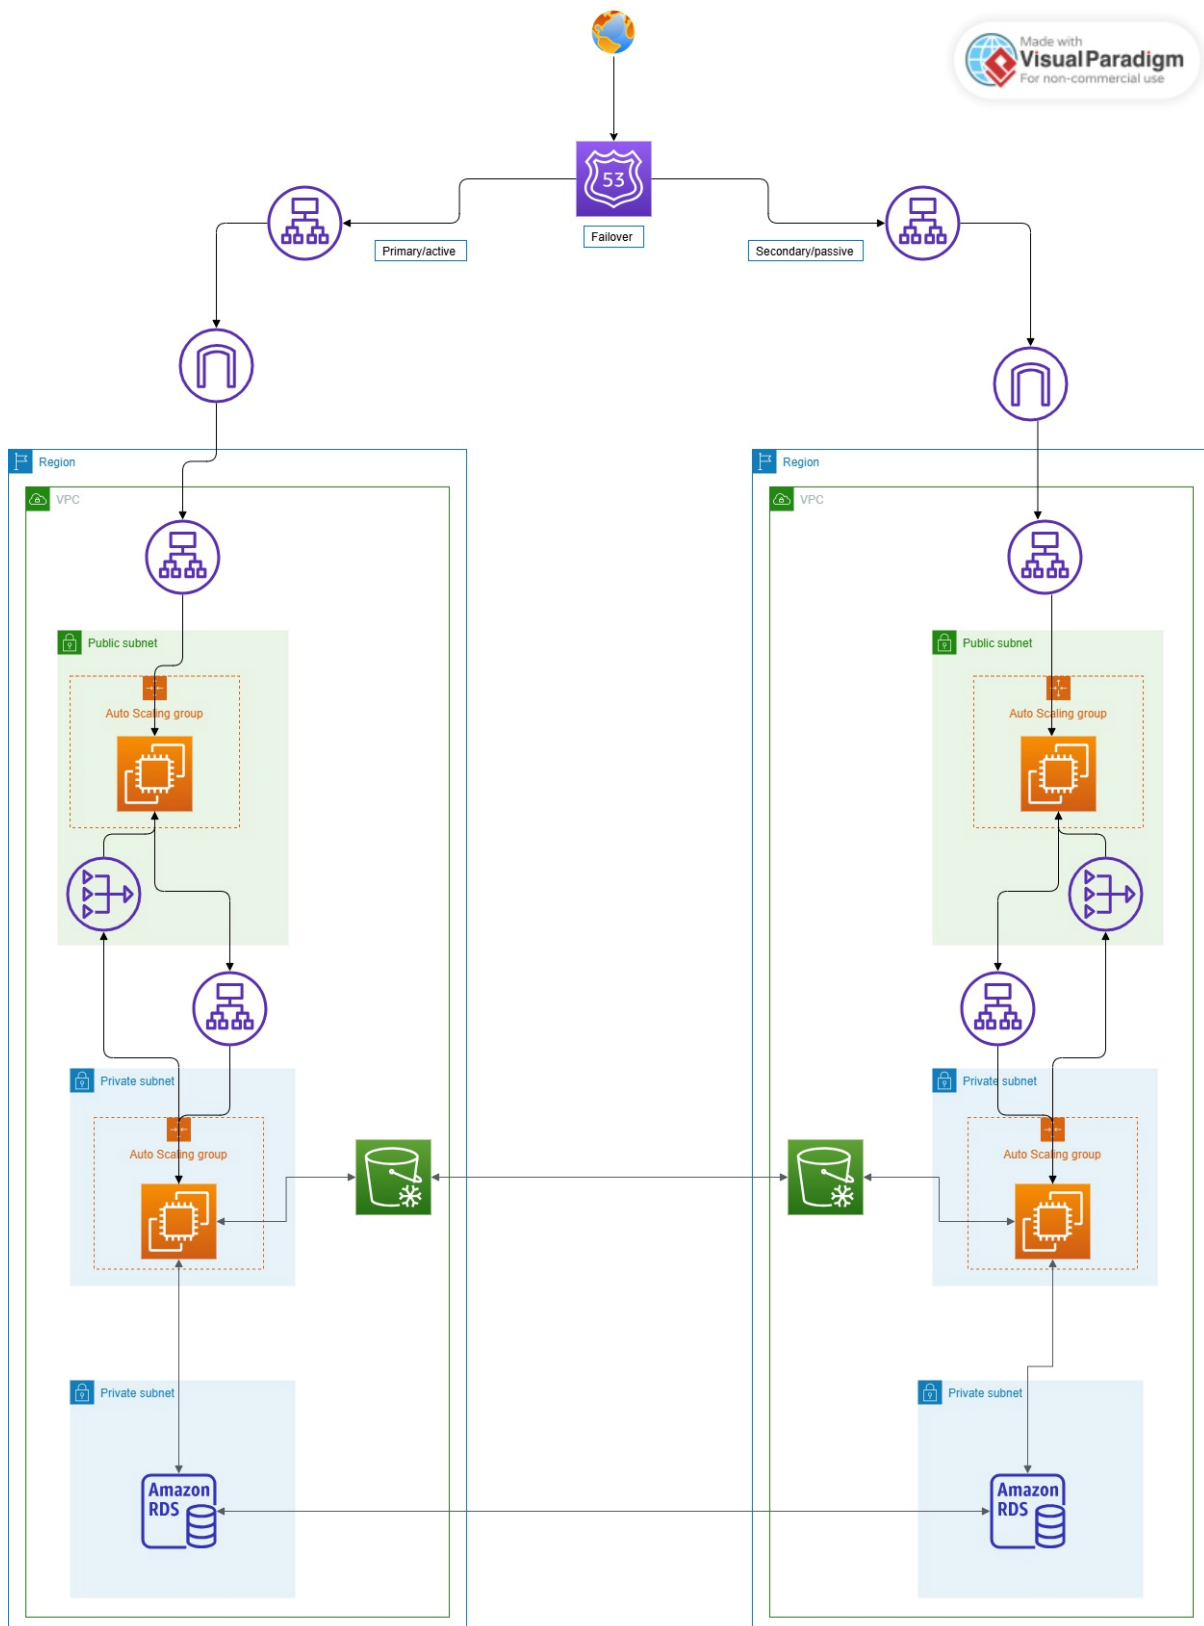
To design and implement a multi-region disaster recovery solution using AWS services, where infrastructure is deployed in two AWS regions. The solution will replicate data across regions, ensure high availability, and implement a failover mechanism for disaster recovery. The project involves using Terraform to provision infrastructure, and CI/CD pipelines to ensure both regions remain synchronized for disaster recovery and failover.

Services:

Component	Role in Architecture
Amazon VPC	Creates isolated networks in each region to host your resources securely.
Amazon S3	Stores data (e.g., app data, backups) and replicates it to another region using S3 Cross-Region Replication (CRR) .
Amazon RDS	Hosts your databases; uses read replica in another region or Aurora Global Database for cross-region sync.
Amazon Route 53	Manages DNS routing; configured with health checks and failover policies so it redirects traffic to the DR region if the primary fails.
Terraform	Automates infrastructure deployment in both regions, ensuring consistency and version control.
Jenkins / AWS CodePipeline	Automates CI/CD workflows to deploy code and infra updates to both regions simultaneously.

Architecture:

To implement a three tier architecture using aws, with implementation of aws services done by terraform script via ci cd pipeline in Jenkins.



Module Structure:

Create a terraform project to launch the aws services. We will be using a modular approach where we create modules for each services and use them in main.tf

DisasterRecovery/

```
main.tf          # Root module orchestrator
variables.tf     # Variable declarations
terraform.tfvars # Environment-specific values
networking/      # VPC, subnets, route tables
security_groups/ # Tiered security group definitions
alb_asg/         # ALB and Auto Scaling Group modules
rds/            # RDS instance configuration
s3/             # Cross-region replicated S3 buckets
```

Architecture Components

Networking & VPC

- Two VPCs (one per region)
- Three subnet tiers per VPC:
 - **Public Subnets:** Host public ALBs and Web EC2 instances
 - **Application Subnets:** Host internal ALBs and App EC2 instances
 - **Database Subnets:** Host RDS in private subnets
- NAT Gateways in public subnets for outbound internet access from private subnets

Security Groups

- **External ALB SG:** Allows inbound HTTP/HTTPS traffic from the internet
- **Web EC2 SG:** Allows inbound HTTP traffic from the ALB
- **App EC2 SG:** Allows inbound traffic from Web tier on port 8080
- **DB SG:** Allows inbound MySQL traffic from the App tier only

Compute & Load Balancing

- **4 ALBs:**
 - 2 Web ALBs (one per region, public)
 - 2 App ALBs (one per region, internal)
- **4 Auto Scaling Groups:**
 - Web and App tiers in each region
- Launch templates install Apache and configure test endpoints for validation

Storage

- RDS MySQL instance deployed in the primary region (us-east-1)
- Cross-region S3 replication for backups and Terraform state sharing

High Availability & Disaster Recovery

- **Route 53 Failover Routing:**
 - Primary ALB (us-east-1) receives traffic when healthy
 - Secondary ALB (us-east-2) automatically takes over if the primary becomes unhealthy
- Multi-region architecture ensures minimal downtime and robust resilience